


# **Pneumonia Prediction from Chest X-ray Images**

(Under the supervision of Prof. Sujoy K. Biswas)

Srijan Mallick  
M.Sc. 1<sup>st</sup> year  
Big Data Analytics  
Reg. No: B2030050

# Aim and Motivation

- Pneumonia is an infection that inflames air sacs in one or both lungs, which may fill with fluid, leading to a difficulty in breathing.
- Life threatening to infants and elderly.
- A deep learning model which would accurately predict pneumonia would aid radiologists and save time and cost.



# Data Overview

- ♦ Data source :  
<https://www.kaggle.com/paultimothymooney/chest-xray-pneumonia>
- ♦ Data comprises of 5856 chest X-ray images, obtained from pediatric patients of 1-5 years of age, from the Guangzhou Women and Children's Medical Center, China.
- ♦ Images divided in 2 classes : Normal and Pneumonia




Normal



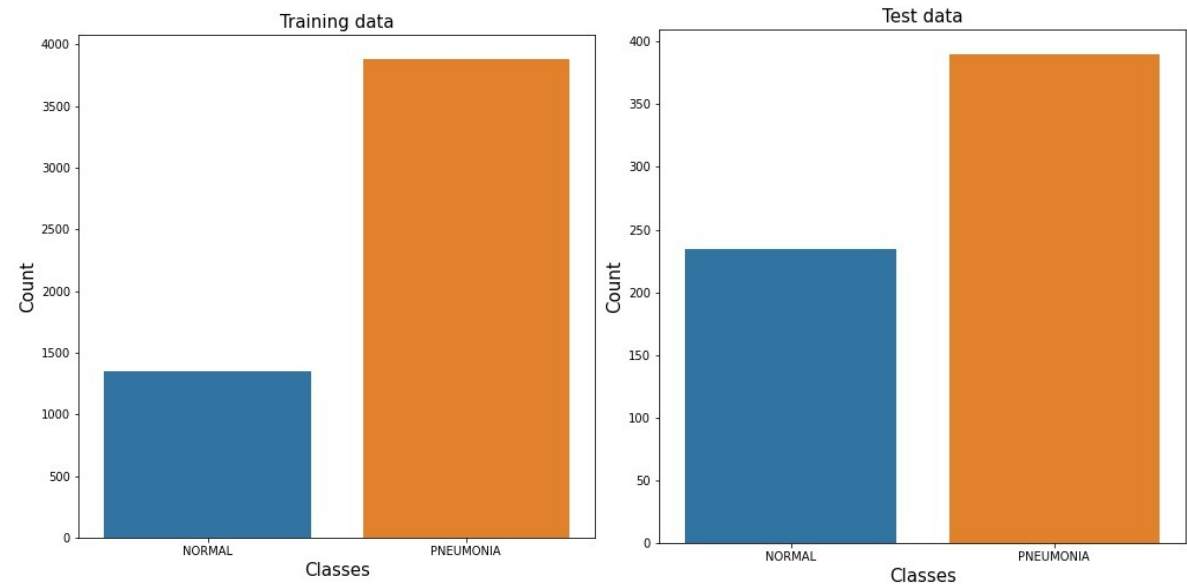
Pneumonia

# Data Visualisation

Final split:      ♦ Train – 4186      ♦ Validation –1046      ♦ Test - 624




Batch of augmented Chest X-ray images of size 20



Data Statistics

# CNN Architecture Overview



```
ConvNet(  
  (conv1): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
  (conv2): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
  (conv3): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
  (conv4): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
  (conv5): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
  (bn1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
  (bn2): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
  (bn3): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
  (bn4): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
  (bn5): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
  (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
  (fc1): Linear(in_features=16384, out_features=1024, bias=True)  
  (fc2): Linear(in_features=1024, out_features=512, bias=True)  
  (fc3): Linear(in_features=512, out_features=2, bias=True)  
  (dropout): Dropout(p=0.5, inplace=False)  
)
```

8-layer deep CNN model used

- Convolution
- Batch Normalisation
- ReLU
- Max Pool
- Dropout


Involved operations

# CNN Architecture Overview

## The need for Batch Normalisation

- **Covariate shift** : Different distribution of the features in different parts of the dataset.
- **Internal Covariate shift** : Covariate shift observed within the network layers due to the distribution of weights and activations.
- **Batch Normalisation** ensures common means and variances of layer inputs which helps in faster convergence and improved gradient flow through the network.

Ensure the output statistics of a layer are fixed.



# CNN Architecture Overview

## Hyperparameters used

- **Adam optimiser** with  
Learning rate : 0.008, with  
Exponential Decay of 0.2  
Weight decay : 0.02  
Epochs : 30
- **SGD optimiser** with  
Learning rate : 0.008, with  
Exponential Decay of 0.2  
Momentum : 0.9  
Nesterov accelerated gradient  
Epochs : 20




(a) SGD without momentum




(b) SGD with momentum

- **Momentum** **blindly accelerates** down slopes: First computes gradient, then makes a big jump.
- Nesterov accelerated gradient (NAG) [Nesterov, 1983] first makes a **big jump** in the direction of the previous accumulated gradient  $\theta - \gamma v_{t-1}$ . Then measures where it ends up and makes a **correction** resulting in the **complete update vector**.

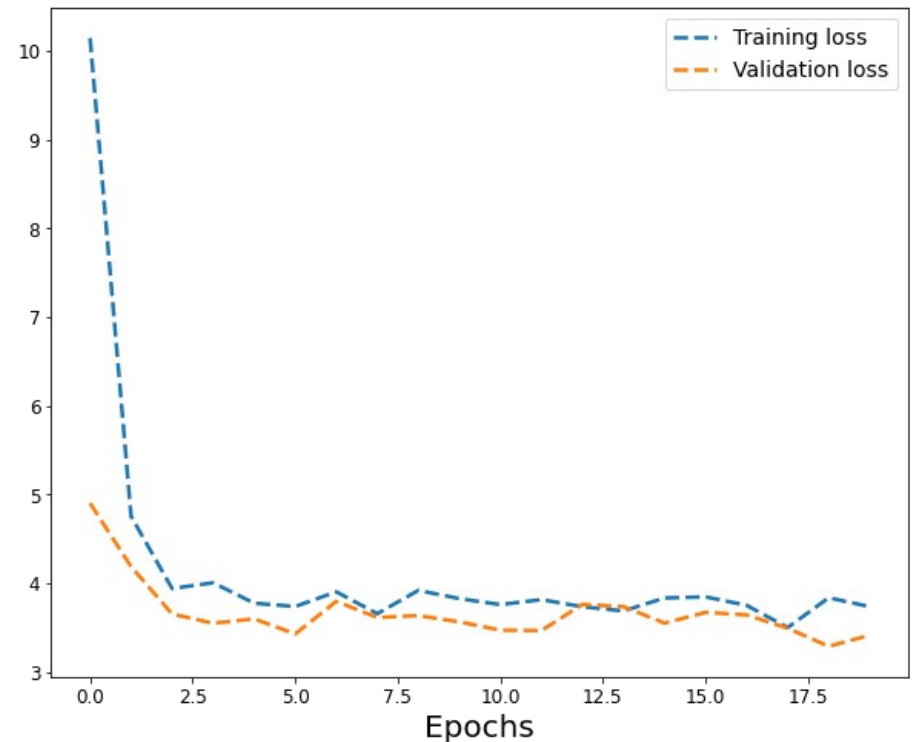
$$v_t = \gamma v_{t-1} + \eta \nabla_{\theta} J(\theta - \gamma v_{t-1})$$
$$\theta = \theta - v_t$$



# Comparing Training/Validation loss




Using Adam optimizer



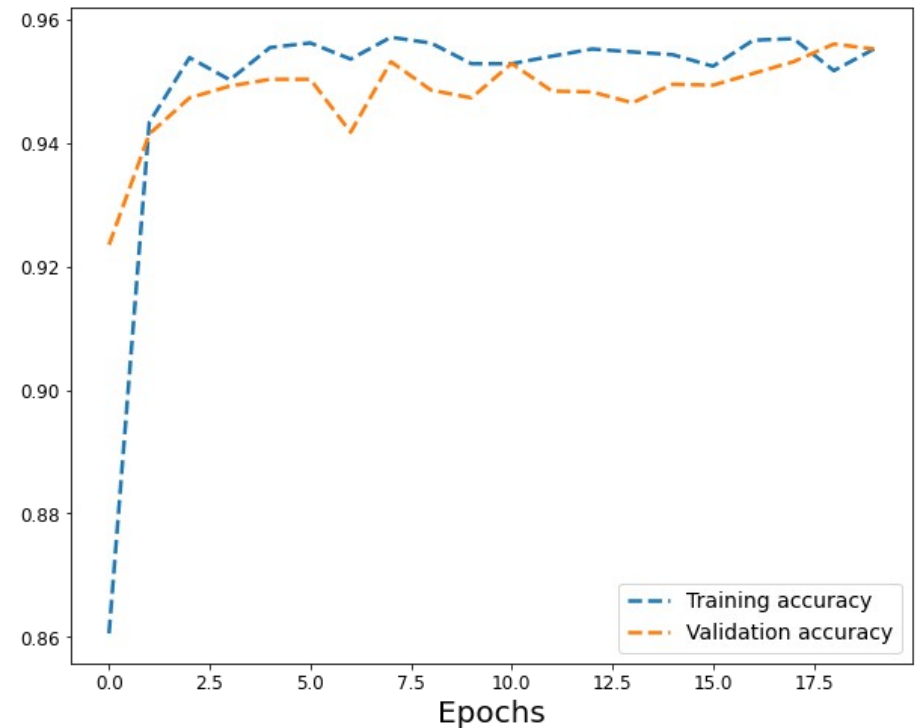
Using SGD with  
Nesterov acceleration



# Comparing Training/Validation accuracy

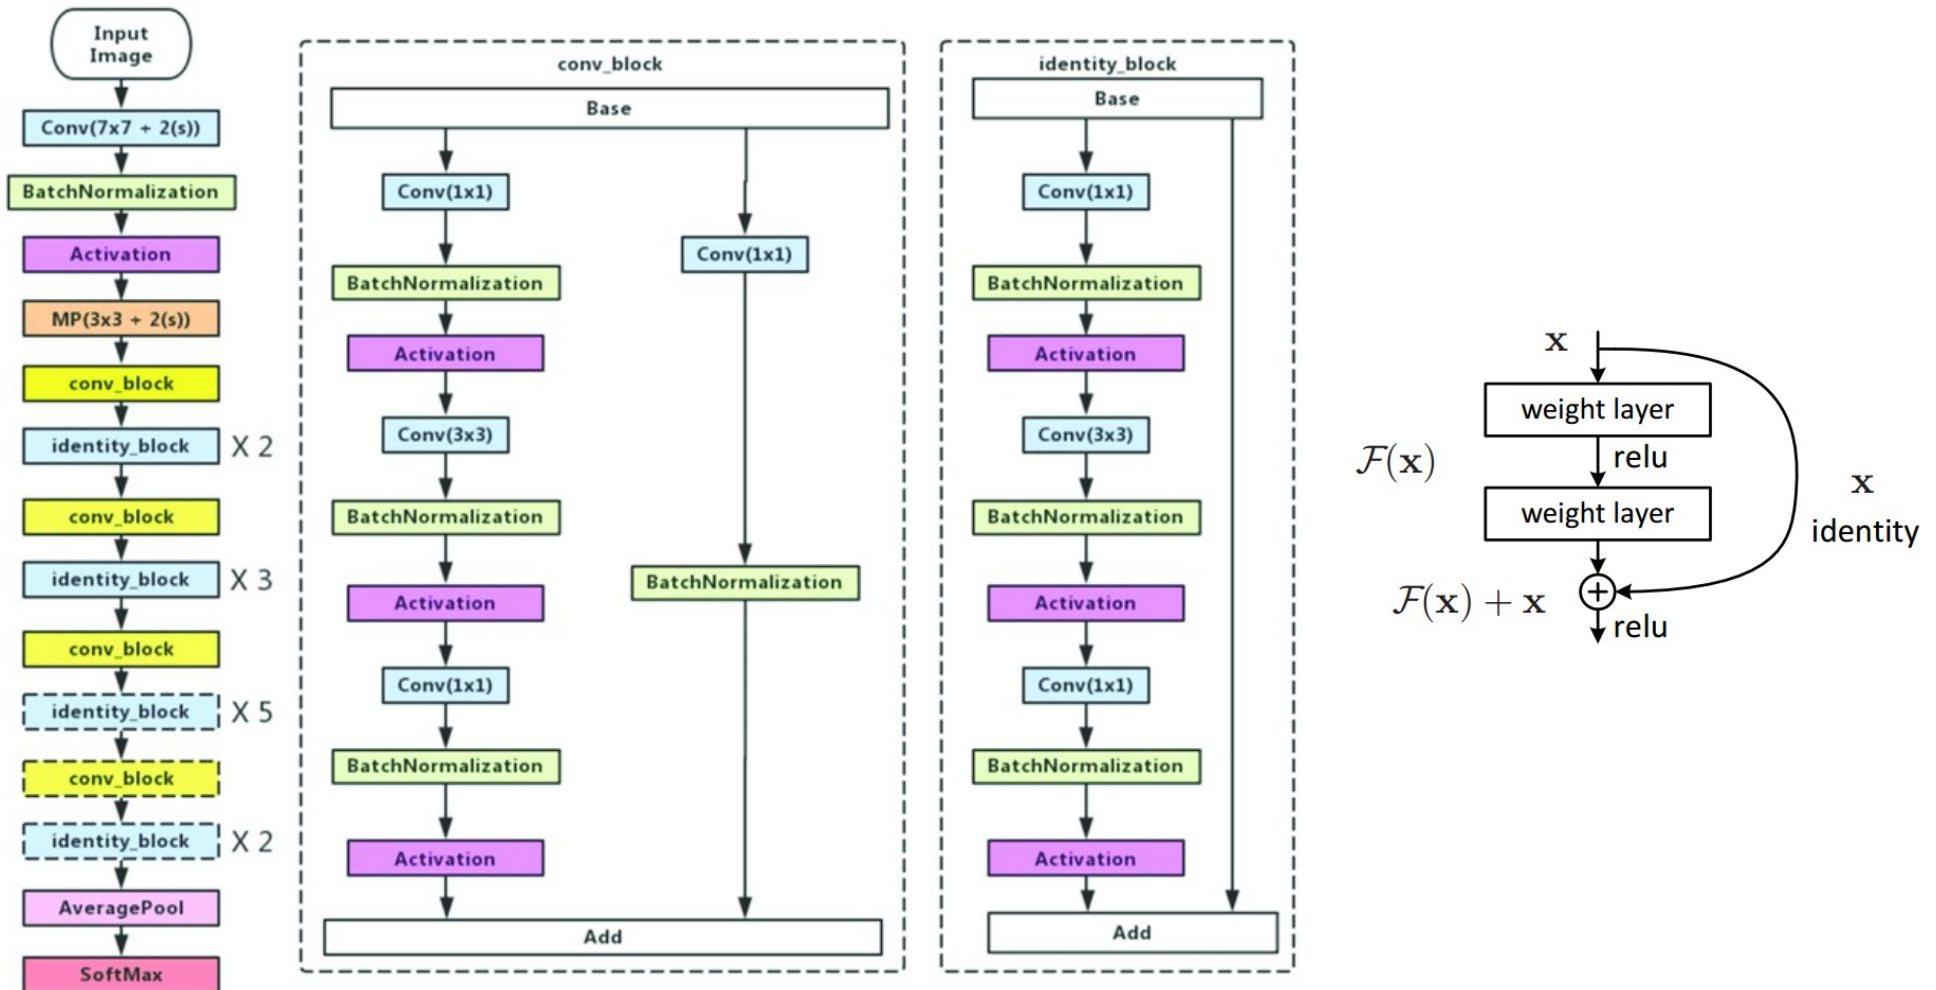


Using Adam optimizer




Using SGD with  
Nesterov acceleration

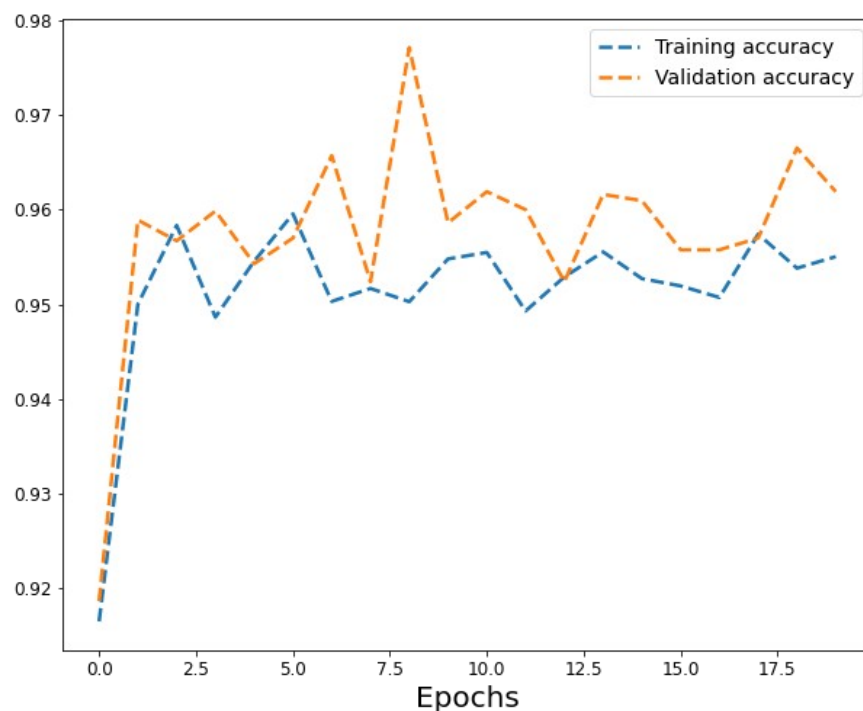
# Transfer Learning using ResNet50 backbone



# Loss and Accuracy obtained using ResNet50 backbone



Training/Validation Losses



Training/Validation Accuracy

# Comparing the different approaches

		8 layer-deep CNN model		ResNet50 model
		Adam	SGD with Nesterov Acceleration	
Min Training loss		4.8844	3.4412	3.2076
Min Validation loss		3.8515	3.0245	2.3778
Test Loss		0.5391	0.7645	0.3771
Highest Training accuracy		94.19%	95.71%	95.95%
Highest Validation accuracy		95.39%	96.36%	97.71%
Test accuracy	NORMAL	49%	49%	73%
	PNEUMONIA	97%	98%	95%
Overall test accuracy		79%	80%	87%


# Visual Saliency

- How does the neural network arrive at its decision ?
- Which portions of an image are ‘salient’, ie play a more important role in this process ?

Saliency maps specifically plot the gradient of the predicted class from the model with respect to the input, or pixel values.

$$Y_c = \text{score of class } c$$
$$\text{saliency} = \max_{r,g,b} \left( \left| \frac{\partial Y_c}{\partial I} \right| \right)$$

The Image and Its Saliency Map




An improved approach : Selvaraju, R.R. et al, “*Grad-CAM: Visual Explanations from Deep Networks via Gradient-based Localization*”, Link : <https://arxiv.org/abs/1610.02391>

# Evaluation Metrics

$$\textit{precision} = \frac{TP}{TP + FP}$$

$$\textit{recall} = \frac{TP}{TP + FN}$$

$$F1 = \frac{2 \times \textit{precision} \times \textit{recall}}{\textit{precision} + \textit{recall}}$$



How many selected items are relevant?

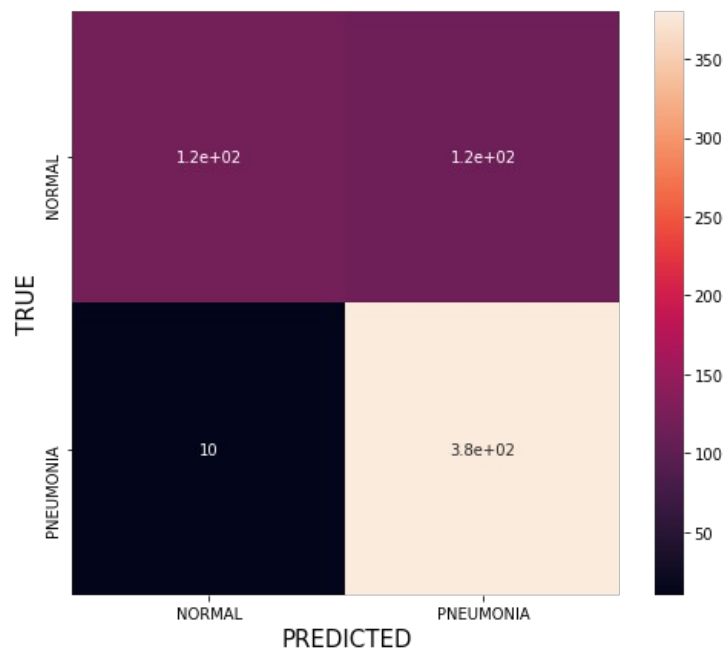
$$\textit{Precision} = \frac{\text{Green}}{\text{Green} + \text{Red}}$$

How many relevant items are selected?

$$\textit{Recall} = \frac{\text{Green}}{\text{Green} + \text{Grey}}$$


# Model Evaluation

Confusion Matrix obtained from  
8-layer deep CNN model



Precision: 0.7677  
Recall: 0.9744  
F1 Score: 0.8588

Confusion Matrix obtained from  
ResNet50 model




Precision: 0.8575  
Recall: 0.9564  
F1 Score: 0.9042

# Prediction Results

PNEUMONIA (PNEUMONIA)




NORMAL (PNEUMONIA)



NORMAL (NORMAL)



NORMAL (NORMAL)



NORMAL (NORMAL)



PNEUMONIA (PNEUMONIA)



PNEUMONIA (PNEUMONIA)



NORMAL (NORMAL)





Thank You !